Stop Job Fraud Without Killing User Trust

The Hybrid Machine Learning that Thinks Like Your Best Human Reviewer

The Human Problem We Solve

Job seekers encounter fraud that wastes their time and erodes trust in digital career platforms. In our 18K job sample, nearly 5% were fraudulent, meaning 1 in 20 applications leads nowhere. Current fraud detection either misses sophisticated scams or flags legitimate posts, frustrating everyone.

Meanwhile, our teams are stuck in tactical fraud-fighting instead of strategic trust-building.

Our Solution: 24/7 ML That Will Never Need a Promotion or Become a Burden on Budgeting Concerns

A hybrid ML that works around the clock without budget negotiations or promotion cycles, freeing your managers from tactical oversight to focus on strategic trust initiatives.

It combines company data checking with language pattern analysis, exactly how humans naturally detect fraud, but at platform scale.

Why This Model Wins

- Best of both worlds: just like human reasoning, check for the company description and the facts.
- Sophisticated fraudsters now complete all company fields correctly but reveal themselves through subtle language patterns like "urgent hiring" combined with vague job duties. Our hybrid model catches this - single-approach systems miss it.
- Performance result: After testing four different ML approaches against 18,000 real job listings, our hybrid model potentially catches 3 out of 4 fraudulent posts.

Here are some examples of patterns that our analysis identified:

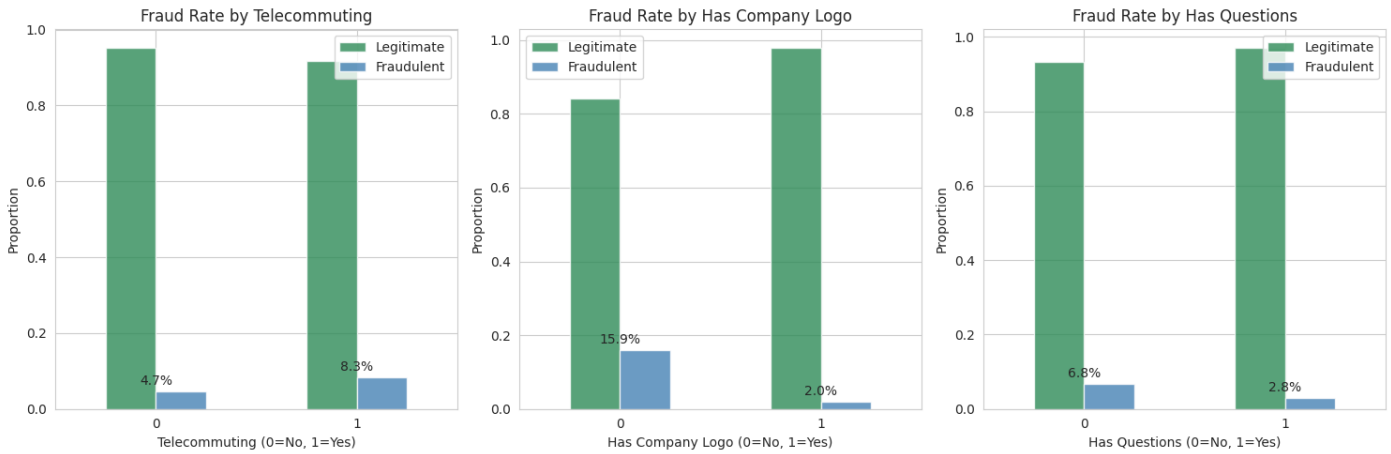

The Behavioural Economics Impact

Trust cascade effect: Reliable fraud prevention increases engagement with legitimate job posts.
Loss aversion addressed: Platform becomes ‘safe by default’ rather than ‘risky until proven otherwise’.

Business translation: Higher user retention, better employer attraction, reduced legal liability.

Implementation Proposal

- Technical deployment: This phase takes 4-6 weeks. During this time, the tool will be integrated, tested, and evaluated. Based on these results, a plan will be made for the roll-out to production.
- Human integration: Automates routine fraud screening, reducing the need for constant supervision. This frees up review teams to focus on edge cases and strategic trust initiatives.
- Resource needs: 1 ML engineer working with 1 senior review member, integration of the existing infrastructures.

Next step: 30-minute technical demo to show the model in action on real fraud cases.